

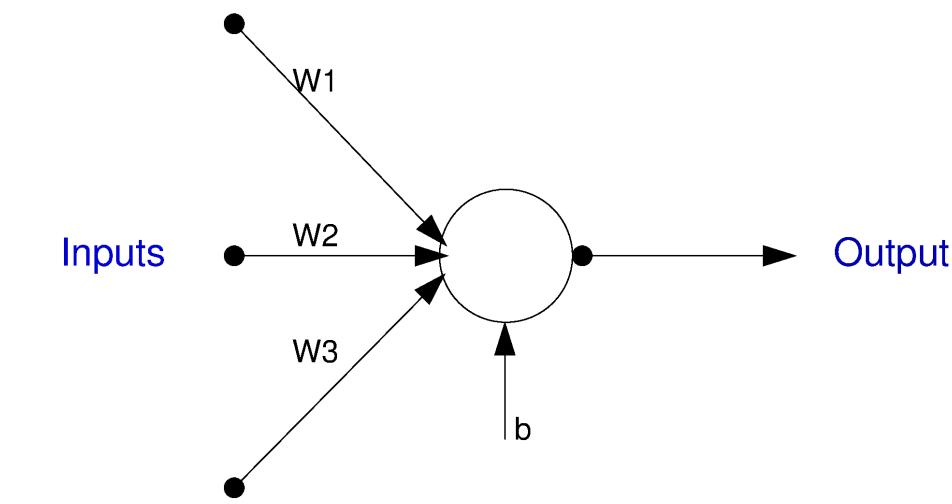
Applied Machine Learning

Training a Simple Neural Network

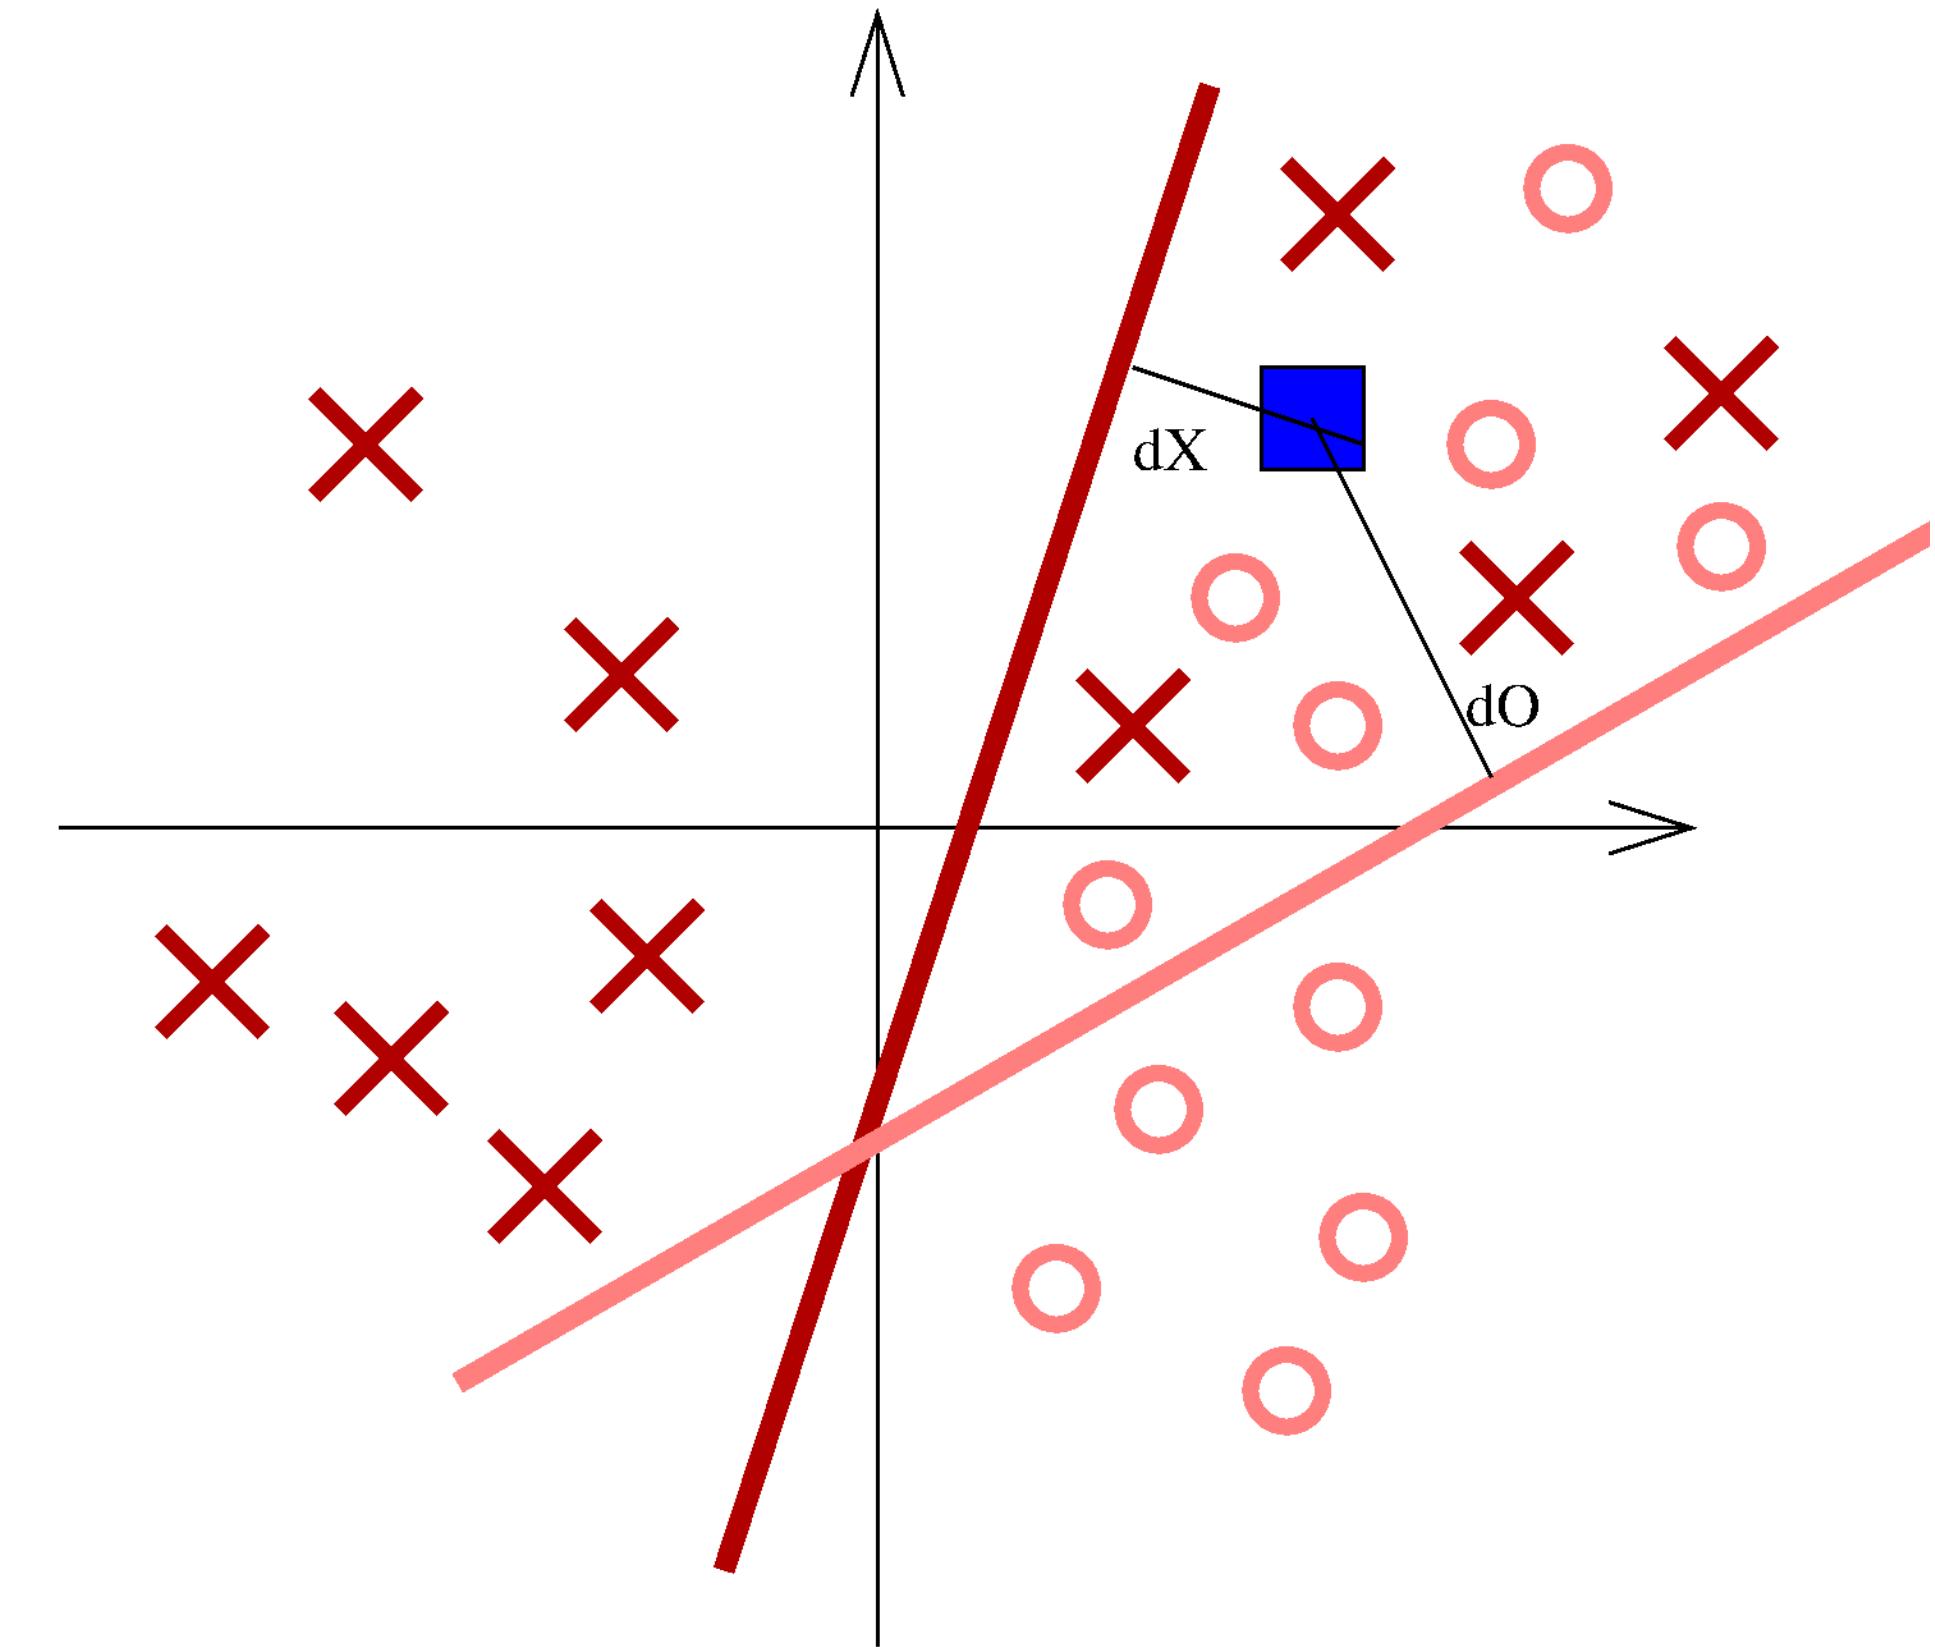
Training a Simple Neural Network

- Cost function for simple neural network
- Gradient
- Stochastic Gradient Descent
- Cross-Validation

Simple Neural Network



- Linear splitting: $u = \sum_i w_i x_i + b$.
- Rectified Linear Unit - **ReLU**: $F(u) = \max(0, u)$
- softmax(u) = $s(\mathbf{o}) = \frac{1}{\sum_k e^{o_k}} \begin{bmatrix} e^{o_1} \\ \vdots \\ e^{o_C} \end{bmatrix}$
- Probability of input \mathbf{x}_i corresponding to class j
 $p(y_j = 1 | \mathbf{x}_i, \mathbf{w}^{(j)}, b^{(j)}) = s_j(\mathbf{o}(\mathbf{x}_i, \mathbf{w}^{(j)}, b^{(j)}))$



Simple Neural Network: Cost function

- Goal: find the vectors of parameters for all the units:

$$\theta^{(j)} = \begin{bmatrix} \mathbf{w}^{(j)} \\ b^{(j)} \end{bmatrix}$$

- Cost function:

- Loss: Maximize log likelihood of training data under probability model

- Penalty: Minimize parameters

- $S(\theta, x; \lambda) = \text{loss} + \lambda \text{ penalty}$

- Loss:

- Minimize $-\log p(\mathbf{y}_i | \mathbf{x}_i, \theta)$

- inputs \mathbf{x}_i , parameters θ , outputs \mathbf{y}_i , N items in dataset

- Log Loss or Cross-Entropy Loss

$$\frac{1}{N} \sum_i [-\log p(\mathbf{y}_i | \mathbf{x}_i, \theta)] =$$

$$\bullet = \frac{1}{N} \sum_{i=1}^N [-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta))]$$

- Penalization: $\frac{1}{2} \sum_{i=1}^C \mathbf{w}^{(i)\top} \mathbf{w}^{(i)}$

- Cost Function: $S(\theta, x; \lambda) =$

$$\bullet = \frac{1}{N} \sum_{i=1}^N [-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta))] + \lambda \frac{1}{2} \sum_{i=1}^C \mathbf{w}^{(i)\top} \mathbf{w}^{(i)}$$

Simple Neural Network: Training

- Goal: find the vectors of parameters θ that minimize the cost function
- Approach:
 - Stochastic Gradient Descent
 - It is hard to find the global minimum
 - reduce the loss over time
 - Minibatches: subset of dataset formed by selecting M items uniformly at random
 - Training
 - Iterate
 - form minibatch with M items from dataset selected uniformly at random
 - Apply Stochastic Gradient Descent to minibatch

Simple Neural Network: Gradient

- Cost Function:
$$S(\theta, x; \lambda) = \frac{1}{N} \sum_{i=1}^N [-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta))] + \lambda \frac{1}{2} \sum_{i=1}^C \mathbf{w}^{(i)\top} \mathbf{w}^{(i)}$$
- Penalty term for unit j :
$$\lambda \frac{1}{2} \mathbf{w}^{(j)\top} \mathbf{w}^{(j)}$$
- Gradient for the penalty term for unit j with respect to parameters
 - with respect to weight w_a in unit j :
$$\frac{\partial}{\partial w_a^j} \lambda \frac{1}{2} \mathbf{w}^{(j)\top} \mathbf{w}^{(j)} = \lambda w_a^j$$
 - with respect to bias in unit j :
$$\frac{\partial}{\partial b^{(j)}} \lambda \frac{1}{2} \mathbf{w}^{(j)\top} \mathbf{w}^{(j)} = 0$$

Simple Neural Network: Gradient

- Cost Function: $S(\theta, x; \lambda) = \frac{1}{N} \sum_{i=1}^N [-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta))] + \lambda \frac{1}{2} \sum_{i=1}^C \mathbf{w}^{(i)\top} \mathbf{w}^{(i)}$
- Indicator function: $\mathbb{I}(\text{condition}) = \begin{cases} 1 & \text{condition=True} \\ 0 & \text{otherwise} \end{cases}$
- loss term for data item i
 - $-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta)) = \sum_{u=1}^C y_u \log s_u(o(\mathbf{x}_i, \theta))$
 - loss term in unit u : $y_u \log s_u(o(\mathbf{x}_i, \theta))$
- Gradient for the loss term for data item i with respect to parameters
 - with respect to weight w_a in unit j
 - $\frac{\partial}{\partial w_a^{(j)}} y_u \log s_u(o(\mathbf{x}_i, \theta)) = y_u \left[\sum_v \frac{\partial \log s_u}{\partial o_v} \frac{\partial o_v}{\partial w_a^{(j)}} \right]$
 - with respect to bias in unit j
 - $\frac{\partial}{\partial b^{(j)}} y_u \log s_u(o(\mathbf{x}_i, \theta)) = y_u \left[\sum_v \frac{\partial \log s_u}{\partial o_v} \frac{\partial o_v}{\partial b^{(j)}} \right]$
- $\frac{\partial \log s_v}{\partial o_v} = \frac{e^{o_v}}{\sum_k e^{o_k}}$
- $\log s_v = \log \left(\frac{e^{o_v}}{\sum_k e^{o_k}} \right) = o_v - \log \sum_k e^{o_k}$
- $\frac{\partial \log s_v}{\partial o_v} = \begin{cases} 1 - \frac{e^{o_v}}{\sum_k e^{o_k}} = 1 - s_v & \text{if } u = v \\ 0 - \frac{e^{o_v}}{\sum_k e^{o_k}} = 0 - s_v & \text{if } u \neq v \end{cases}$
- $\frac{\partial \log s_u}{\partial o_v} = \mathbb{I}_{u=v} - s_v$

Simple Neural Network: Gradient

- Cost Function: $S(\theta, x; \lambda) =$
 - $= \frac{1}{N} \sum_{i=1}^N [-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta))] + \lambda \frac{1}{2} \sum_{i=1}^C \mathbf{w}^{(i)\top} \mathbf{w}^{(i)}$
- loss term for data item i
 - $-\mathbf{y}_i \log \mathbf{s}(\mathbf{o}(\mathbf{x}_i, \theta)) = -\sum_{u=1}^C y_u \log s_u(o(\mathbf{x}_i, \theta))$
 - loss term in unit u : $y_u \log s_u(o(\mathbf{x}_i, \theta))$
- Gradient for the loss term for data item i with respect to parameters
 - with respect to weight w_a in unit j
 - $\frac{\partial}{\partial w_a^{(j)}} y_u \log s_u(o(\mathbf{x}_i, \theta)) = y_u \left[\sum_v \frac{\partial \log s_u}{\partial o_v} \frac{\partial o_v}{\partial w_a^{(j)}} \right]$
 - with respect to bias in unit j
 - $\frac{\partial}{\partial b^{(j)}} y_u \log s_u(o(\mathbf{x}_i, \theta)) = y_u \left[\sum_v \frac{\partial \log s_u}{\partial o_v} \frac{\partial o_v}{\partial b^{(j)}} \right]$

- Indicator function:

$$\mathbb{I}_{\text{condition}} = \begin{cases} 1 & \text{condition=True} \\ 0 & \text{otherwise} \end{cases}$$

- $\frac{\partial \log s_u}{\partial o_v} = \mathbb{I}_{u=v} - s_v$
- $\frac{\partial o_v}{\partial w_a^{(j)}} = x_a^{(j)} \mathbb{I}_{o_v > 0} \mathbb{I}_{u=j}$
- $\frac{\partial o_v}{\partial b^{(j)}} = \mathbb{I}_{o_v > 0} \mathbb{I}_{u=j}$

Simple Neural Network: Training

- Iterate
 - form minibatch with M items from dataset selected uniformly at random
 - Apply Stochastic Gradient Descent with data from minibatch
 - Iteration n
 - $\theta^{(n+1)} = \theta^{(n)} - \eta^{(n)} \nabla_{\theta} \text{cost}$
 - $\eta^{(n)}$: learning rate or step size
 - $e(n)$: epoch for step n
 - Options:
 - $\eta^{(n)} = \frac{1}{a + b \cdot e(n)}$
 - $\eta^{(n)} = \eta(\gamma)^{-e(n)}$ $\gamma > 1$
 - Stopping criteria: based on number of epochs
 - Regularization constant λ : cross-validation
 - Try values of λ at different scales through cross-validation
 - Iterate
 - split dataset into:
 - held-out validation set and training set
 - train network with training set and selected λ
 - evaluate on held-out validation set
 - error is average of held-out errors
 - select λ with smallest held-out error
 - train using the whole dataset

Training a Simple Neural Network

- Cost function for simple neural network
- Gradient
- Stochastic Gradient Descent
- Cross-Validation

Applied Machine Learning

Training a Simple Neural Network