

Applied Machine Learning

Deep Neural Networks and Backpropagation

Deep Neural Networks and Backpropagation

- Backpropagation Algorithm
- Training Deep Neural Networks

Backpropagation Algorithm

- Layer $i - 1$: changes in Loss at output with respect to $\theta^{(i-1)}$
- $\nabla_{\theta^{(i)}} L = \nabla_{\mathbf{o}^{(D)}} L \times \mathbf{J}_{\mathbf{o}^{(D)}; \mathbf{u}^{(D)}} \times \dots \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}}$
- $\mathbf{v}^{(D)} = \nabla_{\mathbf{o}^{(D)}} L$
- $\nabla_{\theta^{(D)}} L = \mathbf{v}^{(D)} \times \mathbf{J}_{\mathbf{o}^{(D)}; \theta^{(D)}}$
- $\mathbf{v}^{(i)} = \mathbf{v}^{(i+1)} \times \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}}$
- $\nabla_{\theta^{(i)}} L = \mathbf{v}^{(i)} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}}$
- Backpropagation
 - Forward pass:
 - $\mathbf{u}^{(1)} = \mathbf{x}$
 - for each layer i from 1 to D:
 - $\mathbf{u}^{(i+1)} = \mathbf{o}^{(i)}(\mathbf{u}^{(i)}, \theta^{(i)})$
 - Backward pass:
 - $\mathbf{v}^{(D)} = \nabla_{\mathbf{o}^{(D)}} L$
 - $\nabla_{\theta^{(D)}} L = \mathbf{v}^{(D)} \times \mathbf{J}_{\mathbf{o}^{(D)}; \theta^{(D)}}$
 - for each layer i from D-1 to 1:
 - $\mathbf{v}^{(i)} = \mathbf{v}^{(i+1)} \times \mathbf{J}_{\mathbf{o}^{(i+1)}; \mathbf{u}^{(i+1)}}$
 - $\nabla_{\theta^{(i)}} L = \mathbf{v}^{(i)} \times \mathbf{J}_{\mathbf{o}^{(i)}; \theta^{(i)}}$

Parameter Initialization

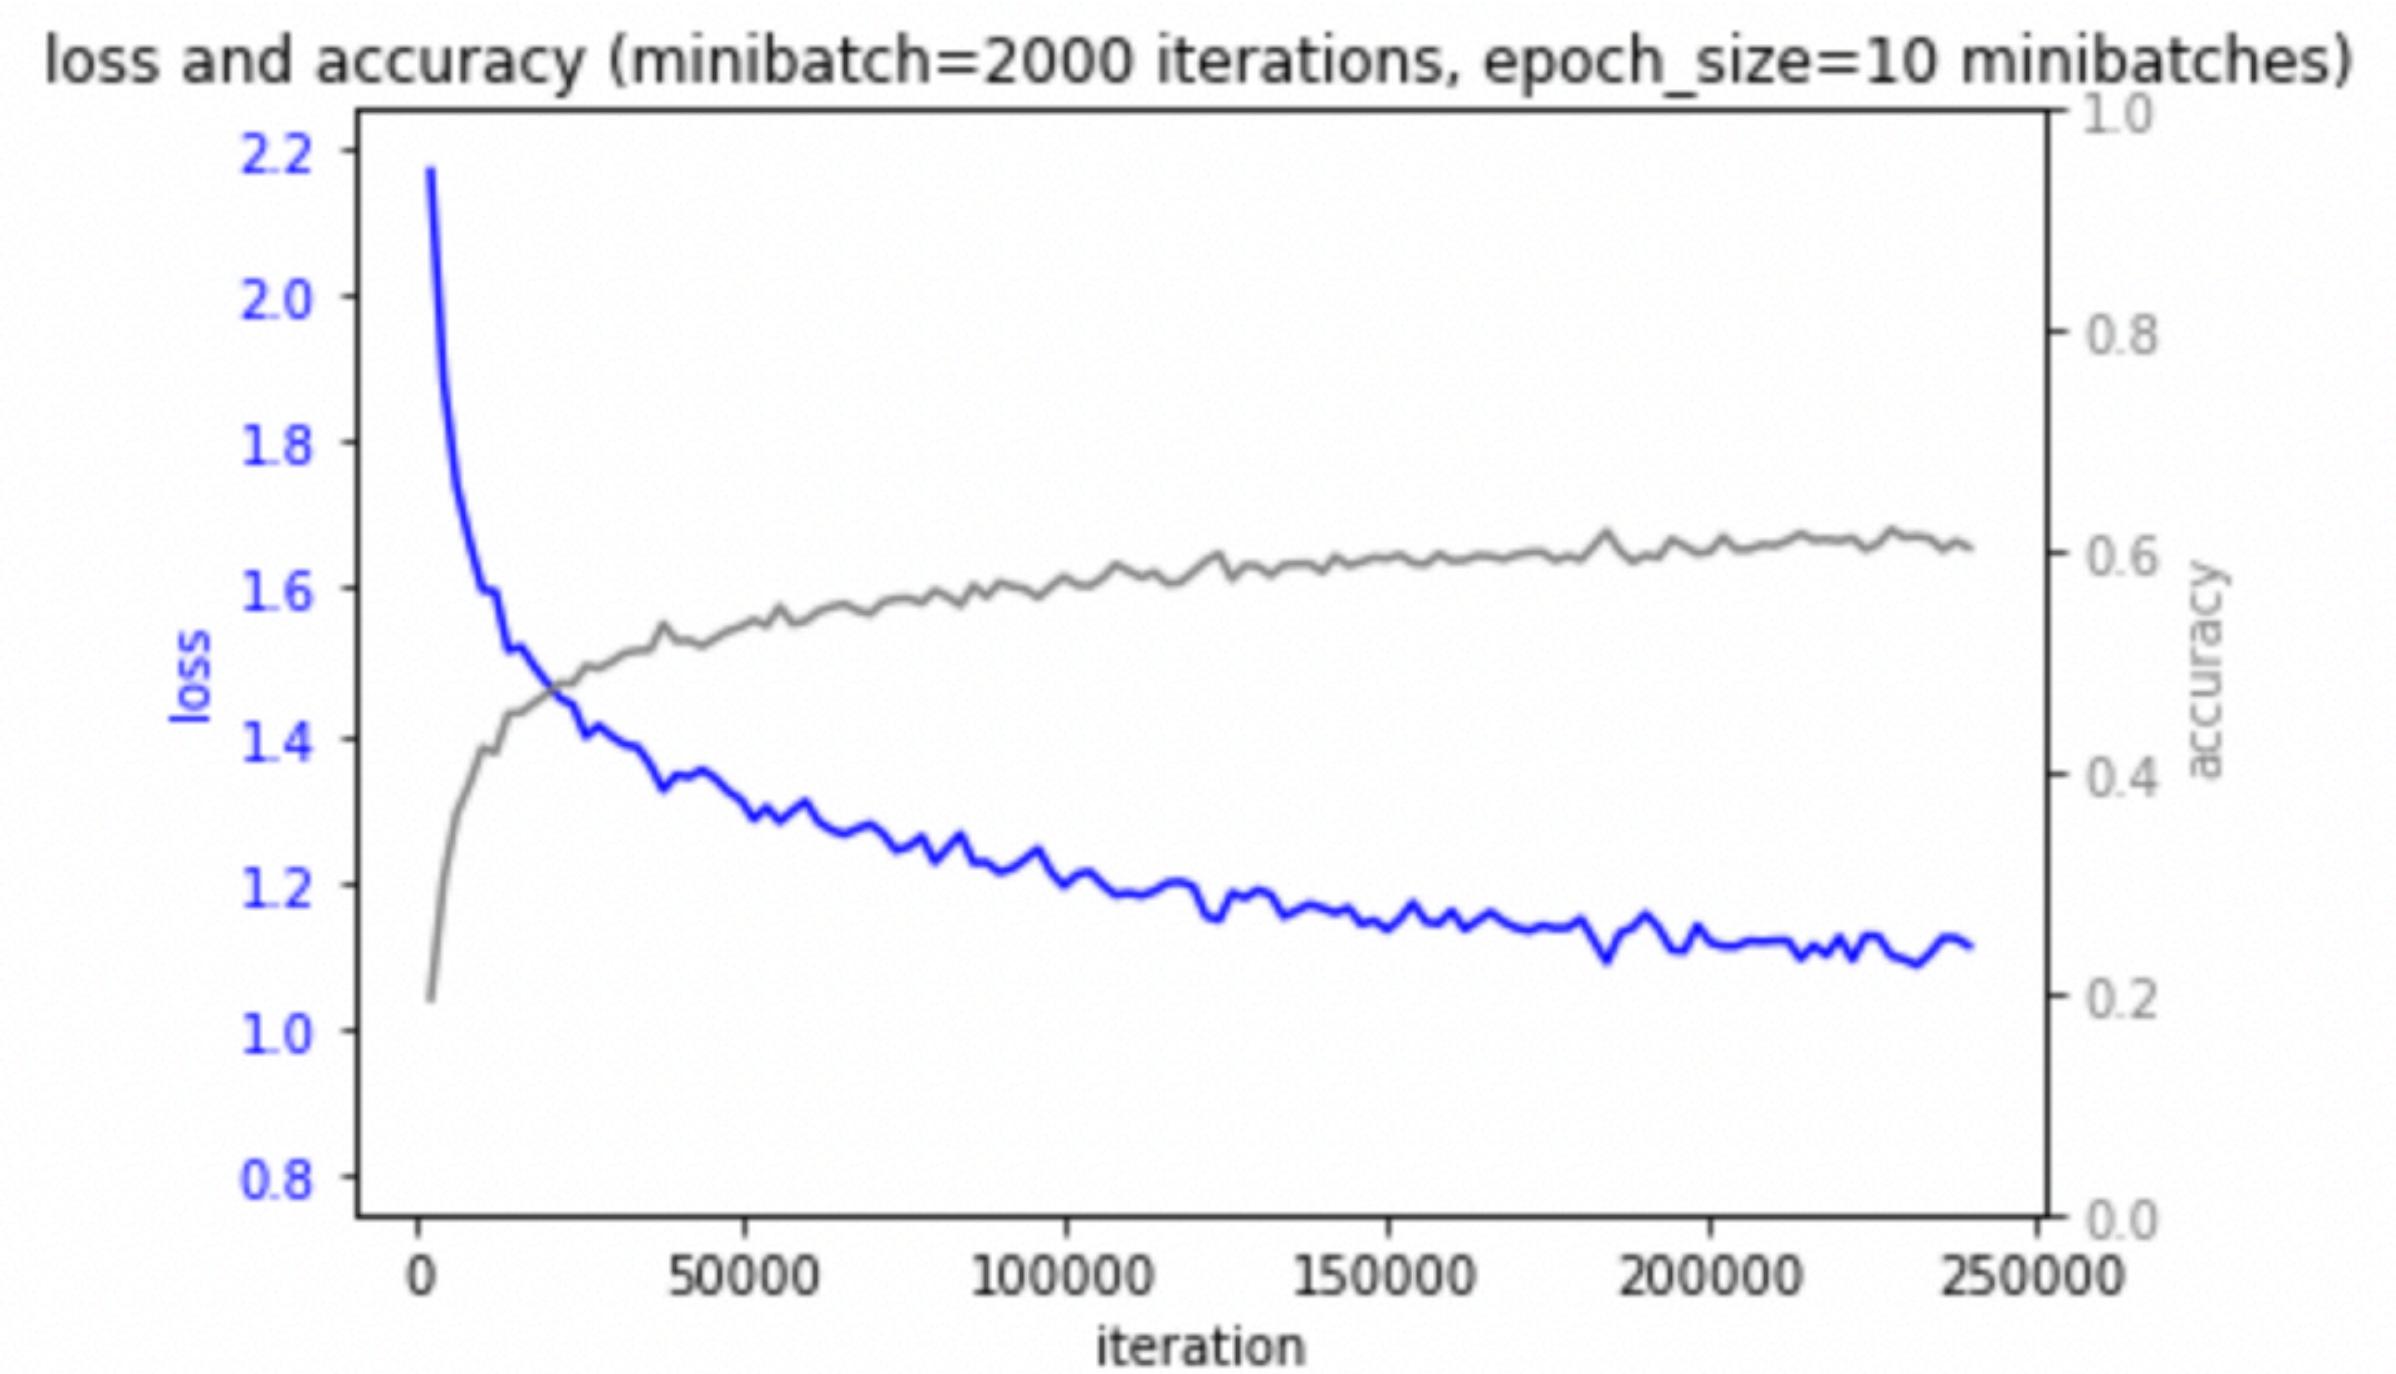
- Initial value of parameters
 - All 0 => bad option
 - Close to their optimal value
 - Each to a random value
 - from zero mean normal distribution with small standard deviation (e.g. 0.01)

Feature Preprocessing

- The relative distribution of different feature inputs may affect learning
 - Large variations in their standard deviations
 - Weights for inputs with larger standard deviations: larger
 - if stepsize is small, it takes longer to build them
 - Weights for inputs with smaller standard deviations: smaller
 - if stepsize is large, it's unlikely to find them
- Preprocessing:
 - Normalize each input to have zero mean and unit standard deviation
 - Apply domain knowledge

Training Loss and Accuracy

- Plots: Loss vs Epochs, Accuracy vs Epochs
- Stepsize or Learning Rate
 - smaller learning rate
 - smoother curves => gradual walk towards minimum
 - more likely to fall into non-optimal loss
 - larger learning rate
 - more noisy curves => jumps on the gradient
 - more likely to diverge
- Reducing loss does not necessarily improve accuracy



Data

- It is hard to successfully train a Deep Neural Network
- Datasets
 - Many times, the more data the better, up to a point

Units

- Units
 - Redundant units
 - May result in cancellation of some at next layer
 - Dropout
 - before each training step
 - set outputs of some randomly selected units to zero
 - do not consider output units
 - Dead units
 - output 0 for every data item => Gradient 0 => dead unit
 - smaller learning rates help
 - many more units

Gradient Obfuscation

- Gradient Obfuscation
 - poor parameter estimates close to output layer => poor gradient update
 - the more layers above, the worst
- initialize with good estimate
- rescaling
- change connectivity structure

Training Deep Neural Networks

- Implementation
 - GPUs
 - APIs
 - describe layers, connectivity, gradients, unit functions
 - map onto the available computing resources (e.g., GPU)
 - training, and evaluation

Deep Neural Networks and Backpropagation

- Backpropagation Algorithm
- Training Deep Neural Networks

Applied Machine Learning

Deep Neural Networks and Backpropagation